

# Chapter 14: Introduction to Boosting

Modern Clinical Data Science  
Chapter Guides  
Bethany Percha, Instructor




# How to Use this Guide

- Read the corresponding notes chapter first
- Try to answer the discussion questions on your own
- Listen to the chapter guide (should be 30 min, max) while following along in the notes

So far...













- Regression
- K-Nearest Neighbors (KNN)
- Decision trees
- Random forests

What are their stories?



## **Ensemble learning:**

Combining multiple models that are only slightly better than random (weak learners) can produce a good model (strong learner).



## **Boosting:**

A principled way of creating weak learners whose combined predictions get better over time.

# AdaBoost

First practical boosting algorithm. For a long time, no one knew why it worked.

1. Initialize the observation weights to  $w_i^{(1)} = \frac{1}{N}$  for  $i = 1, \dots, N$ .

2. For  $m = 1, \dots, M$ :

(a) Select a classifier,  $G_m(x)$ , that minimizes the weighted training error according to the current set of weights,  $w_i^{(m)}$ . Depending on the algorithm, it may be possible to train a single classifier on the weighted training set; in other cases, one may need to select the best-performing classifier from among a predefined set.

(b) Compute

$$\text{err}_m = \frac{\sum_{i=1}^N w_i^{(m)} \cdot \mathcal{I}(y^{(i)} \neq G_m(x^{(i)}))}{\sum_{i=1}^N w_i^{(m)}}$$

(c) Compute voting weight for classifier  $m$ :

$$\alpha_m = \log \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$$

(d) Set

$$w_i^{(m+1)} := w_i^{(m)} \cdot \exp \left[ \alpha_m \cdot \mathcal{I}(y^{(i)} \neq G_m(x^{(i)})) \right]$$

for  $i = 1, \dots, N$ .

3. Output

$$G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$$

Subject ID	friends ( $X_1$ )	money ( $X_2$ )	free time ( $X_3$ )	pet ( $X_4$ )	happy ( $Y$ )
1	1	1	0	0	-1
2	1	1	1	0	-1
3	0	1	1	0	-1
4	0	0	0	0	-1
5	1	0	0	0	-1
6	0	0	0	0	-1
7	1	2	1	0	1
8	1	0	1	0	1
9	0	0	1	1	1
10	1	0	0	1	1

$$X_1 = \begin{cases} 0 & \text{no friends} \\ 1 & \text{friends} \end{cases}$$

$$X_2 = \begin{cases} 0 & \text{poor} \\ 1 & \text{enough money} \\ 2 & \text{rich} \end{cases}$$

$$X_3 = \begin{cases} 0 & \text{no free time} \\ 1 & \text{some free time} \end{cases}$$

$$X_4 = \begin{cases} 0 & \text{no pet} \\ 1 & \text{has a pet} \end{cases}$$



Now, let's go through the process of applying AdaBoost to this dataset, step by step.

- (a) Initialize the observation weights for the training data. (Make them uniform.)

$$w_1^{(1)} =$$

$$w_2^{(1)} =$$

$$w_3^{(1)} =$$

$$w_4^{(1)} =$$

$$w_5^{(1)} =$$

$$w_6^{(1)} =$$

$$w_7^{(1)} =$$

$$w_8^{(1)} =$$

$$w_9^{(1)} =$$

$$w_{10}^{(1)} =$$


Now, let's go through the process of applying AdaBoost to this dataset, step by step.

- (a) Initialize the observation weights for the training data. (Make them uniform.)

$$\begin{aligned}w_1^{(1)} &= \mathbf{0.1} \\w_2^{(1)} &= \mathbf{0.1} \\w_3^{(1)} &= \mathbf{0.1} \\w_4^{(1)} &= \mathbf{0.1} \\w_5^{(1)} &= \mathbf{0.1}\end{aligned}$$

$$\begin{aligned}w_6^{(1)} &= \mathbf{0.1} \\w_7^{(1)} &= \mathbf{0.1} \\w_8^{(1)} &= \mathbf{0.1} \\w_9^{(1)} &= \mathbf{0.1} \\w_{10}^{(1)} &= \mathbf{0.1}\end{aligned}$$

- (b) We will now grow a decision tree on this dataset,  $G_1(x)$ , using these weights. We'll use the `rpart` package in R, just as in Section 7.3. Here is the first tree.




Datapoint ID	$w_i^{(1)}$	happy ( $Y$ )	$G_1(x)$
1	0.1	-1	-1
2	0.1	-1	-1
3	0.1	-1	-1
4	0.1	-1	-1
5	0.1	-1	-1
6	0.1	-1	-1
7	0.1	1	-1
8	0.1	1	-1
9	0.1	1	1
10	0.1	1	1

Compute the misclassification error of this tree,  $\text{err}_1$ . Compare this tree to the one we constructed by hand in Chapter 7.

$$\text{err}_1 =$$

- (b) We will now grow a decision tree on this dataset,  $G_1(x)$ , using these weights. We'll use the `rpart` package in R, just as in Section 7.3. Here is the first tree.



Datapoint ID	$w_i^{(1)}$	happy ( $Y$ )	$G_1(x)$
1	0.1	-1	-1
2	0.1	-1	-1
3	0.1	-1	-1
4	0.1	-1	-1
5	0.1	-1	-1
6	0.1	-1	-1
7	0.1	1	-1
8	0.1	1	-1
9	0.1	1	1
10	0.1	1	1

Compute the misclassification error of this tree,  $\text{err}_1$ . Compare this tree to the one we constructed by hand in Chapter 7.

$$\text{err}_1 = \frac{0.1(2)}{0.1(10)} = 0.2$$

(c) Based on how  $G_1(x)$  performs, calculate  $\alpha_1$ , its voting weight.

$$\alpha_1 =$$

(d) Re-weight the observation weights for the training examples.

$$w_1^{(2)} =$$

$$w_2^{(2)} =$$

$$w_3^{(2)} =$$

$$w_4^{(2)} =$$

$$w_5^{(2)} =$$

$$w_6^{(2)} =$$

$$w_7^{(2)} =$$

$$w_8^{(2)} =$$

$$w_9^{(2)} =$$

$$w_{10}^{(2)} =$$

(c) Based on how  $G_1(x)$  performs, calculate  $\alpha_1$ , its voting weight.

$$\alpha_1 = \log \left( \frac{1 - \text{err}_1}{\text{err}_1} \right) = 1.386$$

(d) Re-weight the observation weights for the training examples.

$$w_1^{(2)} = 0.1$$

$$w_2^{(2)} = 0.1$$

$$w_3^{(2)} = 0.1$$

$$w_4^{(2)} = 0.1$$

$$w_5^{(2)} = 0.1$$

$$w_6^{(2)} = 0.1$$


$$w_7^{(2)} = 0.1 * \exp(1.386) = 0.4$$

$$w_8^{(2)} = 0.1 * \exp(1.386) = 0.4$$

$$w_9^{(2)} = 0.1$$

$$w_{10}^{(2)} = 0.1$$

(e) Now we grow another decision tree,  $G_2(x)$ , using these new weights as inputs to the `rpart` package.




Datapoint ID	$w_i^{(2)}$	happy ( $Y$ )	$G_2(x)$
1	0.1	-1	-1
2	0.1	-1	-1
3	0.1	-1	-1
4	0.1	-1	-1
5	0.1	-1	-1
6	0.1	-1	-1
7	0.4	1	1
8	0.4	1	1
9	0.1	1	1
10	0.1	1	-1

Compute the misclassification error of this tree,  $\text{err}_2$ .

$\text{err}_2 =$

(e) Now we grow another decision tree,  $G_2(x)$ , using these new weights as inputs to the `rpart` package.



Datapoint ID	$w_i^{(2)}$	happy (Y)	$G_2(x)$
1	0.1	-1	-1
2	0.1	-1	-1
3	0.1	-1	-1
4	0.1	-1	-1
5	0.1	-1	-1
6	0.1	-1	-1
7	0.4	1	1
8	0.4	1	1
9	0.1	1	1
10	0.1	1	-1

Compute the misclassification error of this tree,  $\text{err}_2$ .

$$\text{err}_2 = \frac{0.1}{0.1(8) + 0.4(2)} = 0.0625$$



(f) Based on how  $G_2(x)$  performs, calculate  $\alpha_2$ , its voting weight.

$$\alpha_2 = \log \left( \frac{1 - \text{err}_2}{\text{err}_2} \right) = 2.708$$

(g) Re-weight the observation weights for the training examples.

$$w_1^{(3)} = 0.1$$

$$w_2^{(3)} = 0.1$$

$$w_3^{(3)} = 0.1$$

$$w_4^{(3)} = 0.1$$

$$w_5^{(3)} = 0.1$$

$$w_6^{(3)} = 0.1$$


$$w_7^{(3)} = 0.4$$

$$w_8^{(3)} = 0.4$$

$$w_9^{(3)} = 0.1$$

$$w_{10}^{(3)} = 0.1 * \exp(2.708) = 1.5$$

- (h) Now we grow another decision tree,  $G_3(x)$ , using these new weights as inputs to the `rpart` package.



Datapoint ID	$w_i^{(3)}$	happy (Y)	$G_3(x)$
1	0.1	-1	-1
2	0.1	-1	-1
3	0.1	-1	-1
4	0.1	-1	-1
5	0.1	-1	1
6	0.1	-1	-1
7	0.4	1	1
8	0.4	1	1
9	0.1	1	-1
10	1.5	1	1

- (i) Compute the misclassification error of this tree,  $\text{err}_3$ .

$$\text{err}_3 = \frac{0.2}{0.1(7) + 0.4(2) + 1.5} = 0.067$$

(j) Based on how  $G_3(x)$  performs, calculate  $\alpha_3$ , its voting weight.

$$\alpha_3 = \log \left( \frac{1 - \text{err}_3}{\text{err}_3} \right) = 2.639$$

(k) Re-weight the observation weights for the training examples.

$$w_1^{(4)} = 0.1$$

$$w_2^{(4)} = 0.1$$

$$w_3^{(4)} = 0.1$$

$$w_4^{(4)} = 0.1$$

$$w_5^{(4)} = 0.1 * \exp(2.639) = 1.4$$

$$w_6^{(4)} = 0.1$$


$$w_7^{(4)} = 0.4$$

$$w_8^{(4)} = 0.4$$

$$w_9^{(4)} = 0.1 * \exp(2.639) = 1.4$$

$$w_{10}^{(4)} = 1.5$$

(l) Now we grow another decision tree,  $G_4(x)$ , using these new weights as inputs to the `rpart` package.



Datapoint ID	$w_i^{(4)}$	happy (Y)	$G_4(x)$
1	0.1	-1	-1
2	0.1	-1	1
3	0.1	-1	1
4	0.1	-1	-1
5	1.4	-1	-1
6	0.1	-1	-1
7	0.4	1	1
8	0.4	1	1
9	1.4	1	1
10	1.5	1	1

(m) Compute the misclassification error of this tree,  $err_4$ .

err<sub>4</sub> =

0.2

0.1(5) + 0.4(2) + 1.4(2) + 1.5

= 0.036

(n) Based on how  $G_4(x)$  performs, calculate  $\alpha_4$ , its voting weight.

$$\alpha_4 = \log \left( \frac{1 - \text{err}_4}{\text{err}_4} \right) = 3.296$$

(o) Output the weighted average of the four classifiers' votes for each training example:

$$G(x) = \text{sign} [\alpha_1 G_1(x) + \alpha_2 G_2(x) + \alpha_3 G_3(x) + \alpha_4 G_4(x)]$$


What is the final training error?

		1.386	2.708	2.639	3.296	
Datapoint ID	happy ( $Y$ )	$G_1(x)$	$G_2(x)$	$G_3(x)$	$G_4(x)$	$G(x)$
1	-1	-1	-1	-1	-1	
2	-1	-1	-1	-1	1	
3	-1	-1	-1	-1	1	
4	-1	-1	-1	-1	-1	
5	-1	-1	-1	1	-1	
6	-1	-1	-1	-1	-1	
7	1	-1	1	1	1	
8	1	-1	1	1	1	
9	1	1	1	-1	1	
10	1	1	-1	1	1	


### Question 14.1

The most difficult thing to understand in all of this is how the updated weights play into the construction of subsequent trees. A clue comes from the dataset percentages shown in the nodes of each tree. For example, trees 1 and 3 actually

Datapoint ID	$w_i^{(1)}$	happy ( $Y$ )	$G_1(x)$
1	0.1	-1	-1
2	0.1	-1	-1
3	0.1	-1	-1
4	0.1	-1	-1
5	0.1	-1	-1
6	0.1	-1	-1
7	0.1	1	-1
8	0.1	1	-1
9	0.1	1	1
10	0.1	1	1



Datapoint ID	$w_i^{(4)}$	happy ( $Y$ )	$G_4(x)$
1	0.1	-1	-1
2	0.1	-1	1
3	0.1	-1	1
4	0.1	-1	-1
5	1.4	-1	-1
6	0.1	-1	-1
7	0.4	1	1
8	0.4	1	1
9	1.4	1	1
10	1.5	1	1



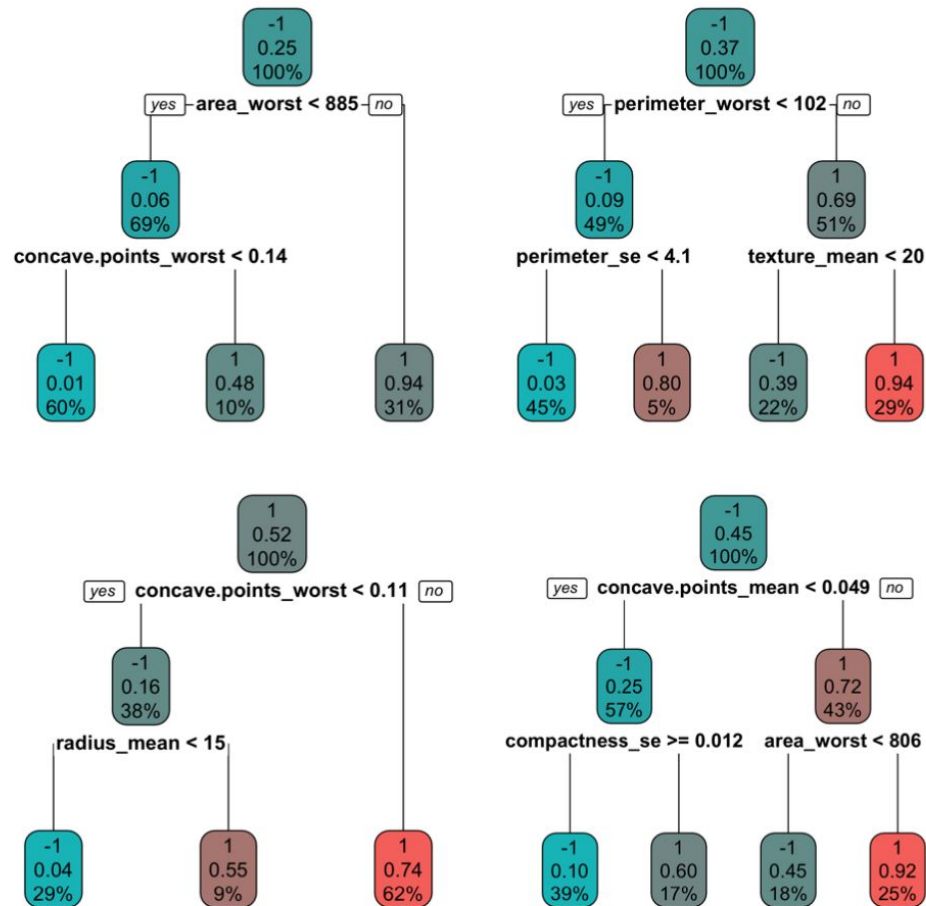

## **Gradient boosting:**

A more general conception of boosting in which later models are fit to the errors (pseudo-residuals) of earlier models.

How does this guarantee diversity?

# Wisconsin Breast Cancer dataset revisited

Ensemble error vs number of trees





### Question 14.2

Compare and contrast boosting and random forests on the basis of:

- Whether each tree uses all or part of the dataset
- Whether they consider a subset of the predictors at each split or all the predictors
- Whether you can parallelize the construction of different trees (i.e., build them at the same time on different processors)
- Whether the votes of different classifiers are independent
- Ease of use and interpretability